

Experiment 1

Aim: Study of ETL process and its tools.

Theory:

Dataset:

	Name	Department	Basic	Bonus
0	Employee_1	HR	41291.0	8181.0
1	Employee_2	Sales	43241.0	8635.0
2	Employee_3	IT	NaN	8896.0
3	Employee_4	HR	47441.0	8685.0
4	Employee_5	Marketing	49767.0	7768.0
...
195	Employee_196	Finance	40048.0	7673.0
196	Employee_197	Finance	49372.0	9106.0
197	Employee_198	Finance	33814.0	6700.0
198	Employee_199	HR	51541.0	9817.0
199	Employee_200	Finance	36934.0	3959.0

Code:

```
import pandas as pd
```

1. Extract: Read CSV

```
employee = pd.read_csv("employee_salary.csv") # upload first in Colab
display(employee.head()) #dipslay initial 5 rows from the table
```

2. Transform

```
# Handle missing values using mean
employee['Basic'] = employee['Basic'].fillna(employee['Basic'].mean())
employee['Bonus'] = employee['Bonus'].fillna(employee['Bonus'].mean())

# Calculate Total_Salary
employee['Total_Salary'] = employee['Basic'] + employee['Bonus']

# Add Tax column
employee['Tax'] = employee['Total_Salary'].apply(lambda x: 0.10 * x if x >= 40000 else 0.05 * x)
```

3. Load: Save as employee_salary_transformed.csv

```
employee.to_csv("employee_salary_transformed.csv", index=False)
```

```
print("Transformation complete! File saved as employee_salary_transformed.csv")
employee
```

Output:

File saved successfully

	Name	Department	Basic	Bonus	Total Salary	Total Tax	
0	Employee_1	HR	41291.0	8181.0	49472.0	4947.2	
1	Employee_2	Sales	43241.0	8635.0	51876.0	5187.6	
2	Employee_3	IT	NaN	8896.0	NaN	NaN	
3	Employee_4	HR	47441.0	8685.0	56126.0	5612.6	
4	Employee_5	Marketing	49767.0	7768.0	57535.0	5753.5	
...	
195	Employee_196	Finance	40048.0	7673.0	47721.0	4772.1	
196	Employee_197	Finance	49372.0	9106.0	58478.0	5847.8	
197	Employee_198	Finance	33814.0	6700.0	40514.0	4051.4	
198	Employee_199	HR	51541.0	9817.0	61358.0	6135.8	
199	Employee_200	Finance	36934.0	3959.0	40893.0	4089.3	

200 rows × 6 columns

Learning Outcomes:

Experiment 2

Aim: Program of Data warehouse cleansing to input names from users (inconsistent) and format them.

Theory:

Dataset:

Code:

```
import pandas as pd
df = pd.read_csv('employees_inconsistent_simple.csv')
print(df.head(5))
```

	ID	Name	Department
0	1	DAvid miLLEr	Finance
1	2	MARyann doE	Operations
2	3	lISa TaYlOr	Sales
3	4	suNiL wILSON	Finance
4	5	SUNiL dAVIS	Finance

```
def cleanse_name(name):
    cleaned = name.strip()
    cleaned = " ".join(cleaned.split())
    cleaned = cleaned.title()
    return cleaned
df['cleaned_name']= df["Name"].apply(cleanse_name) #apply transformation
print(df)
```

Output:

	ID	Name	Department	cleaned_name
0	1	DAvid miLLEr	Finance	David Miller
1	2	MARyann doE	Operations	Maryann Doe
2	3	lISa TaYlOr	Sales	Lisa Taylor
3	4	suNiL wILSON	Finance	Sunil Wilson
4	5	SUNiL dAVIS	Finance	Sunil Davis
..
195	196	anANDA TaylOr	Marketing	Ananda Taylor
196	197	suniL JONES	Finance	Sunil Jones
197	198	rAhul wilSON	Finance	Rahul Wilson
198	199	PeTer kUMAr	HR	Peter Kumar
199	200	maRYanN sMITh	IT	Maryann Smith

[200 rows x 4 columns]

Learning Outcomes:

Experiment 3

Aim: Program of Data warehouse cleansing to remove redundancy in data.

Theory:

Dataset:

Code:

```
from fuzzywuzzy import fuzz
import pandas as pd
import matplotlib.pyplot as plt
#Load csv
df = pd.read_csv("customer_data.csv")
print("Original Dataset:\n", df)
```

Original Dataset:						
	Customer_ID	Name	Email	Phone	Address	
0	101	John Doe	john@example.com	1234567890	123 Elm St	
1	102	Jane Smith	jane@example.com	2345678901	456 Oak St	
2	103	john doe	john@example.com	1234567890	123 Elm Street	
3	104	Alice Brown	alice@example.com	3456789012	789 Pine St	
4	105	Bob Martin	bob@example.com	4567890123	321 Maple Ave	
5	106	Jane Smith	jane@example.com	2345678901	456 Oak Street	
6	107	Jon Doe	jon@example.com	1234567890	123 Elm St	
7	108	Ally Brown	alice@example.com	3456789012	789 Pine Street	

```
#Standardize names and addresses
df['Name'] = df['Name'].str.strip().str.title()
df['Address'] = df['Address'].str.strip().str.title()
print("\nAfter Standardization\n", df)
#Remove exact duplicates
df_no_duplicates = df.drop_duplicates()
print("After removing duplicates\n", df_no_duplicates)
```

After removing duplicates						
	Customer_ID	Name	Email	Phone	Address	
0	101	John Doe	john@example.com	1234567890	123 Elm St	
1	102	Jane Smith	jane@example.com	2345678901	456 Oak St	
2	103	John Doe	john@example.com	1234567890	123 Elm Street	
3	104	Alice Brown	alice@example.com	3456789012	789 Pine St	
4	105	Bob Martin	bob@example.com	4567890123	321 Maple Ave	
5	106	Jane Smith	jane@example.com	2345678901	456 Oak Street	
6	107	Jon Doe	jon@example.com	1234567890	123 Elm St	
7	108	Ally Brown	alice@example.com	3456789012	789 Pine Street	

```
#Detect near Duplicates using fuzzy matching
names = df_no_duplicates['Name'].tolist()
print("Potential Near Duplicates");
for i in range(len(names)):
    for j in range(i + 1, len(names)):
        similarity = fuzz.ratio(names[i],
        names[j]); if similarity > 80: # Threshold
            for near
        duplicates
```

```
        print(f"'{names[i]}' and '{names[j]}'\n(Similarity: {similarity}%)")
```

→ Potential Near Duplicates

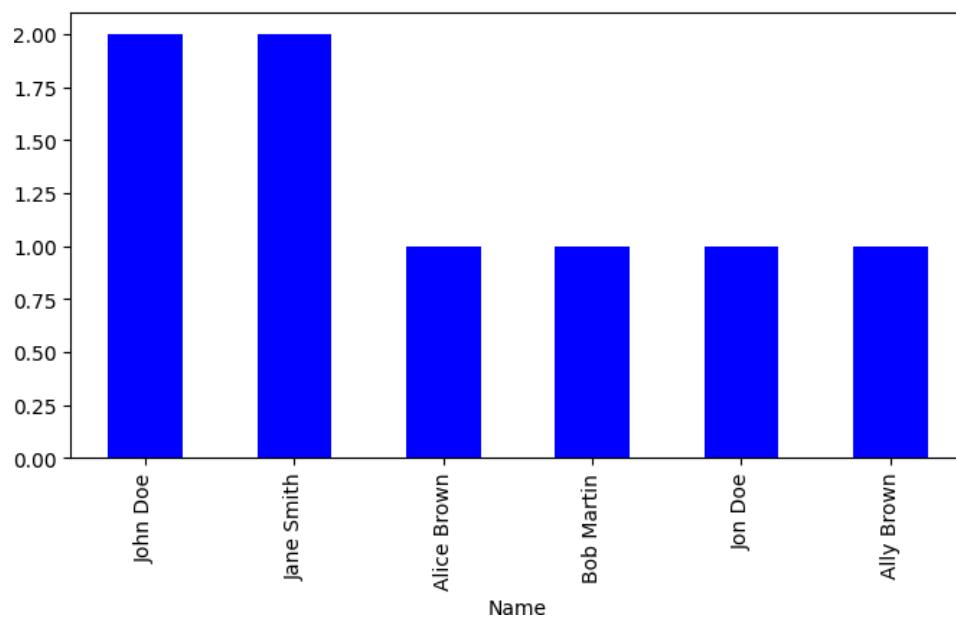
```
'John Doe' and 'John Doe' (Similarity: 100%)
'John Doe' and 'Jon Doe' (Similarity: 93%)
'Jane Smith' and 'Jane Smith' (Similarity: 100%)
'John Doe' and 'Jon Doe' (Similarity: 93%)
```

```
df_no_duplicates=df_no_duplicates.drop_duplicates(subset=['Phone'], keep='first')
print("After Removing Duplicates : \n", df_no_duplicates)
```

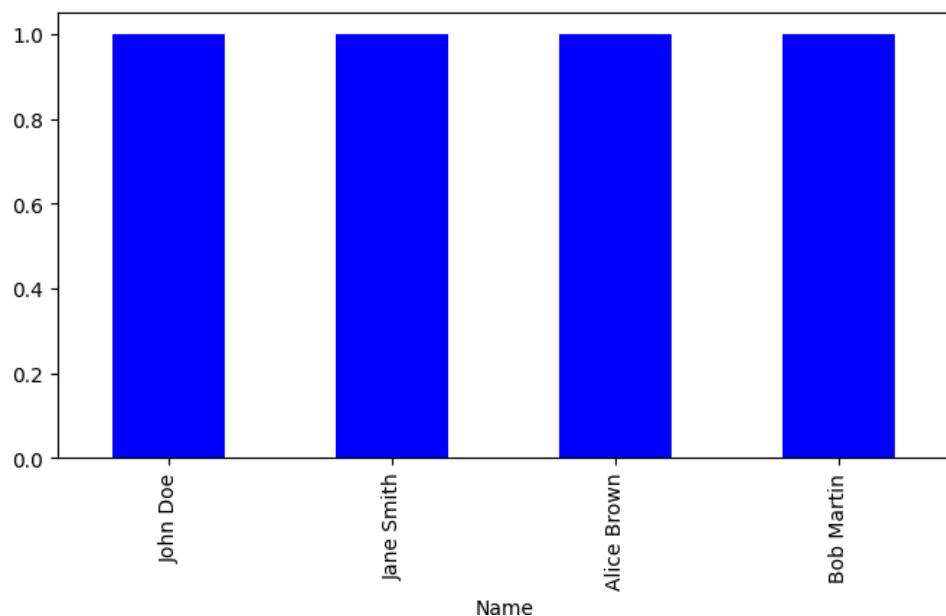
→ After Removing Duplicates :

	Customer_ID	Name	Email	Phone	Address
0	101	John Doe	john@example.com	1234567890	123 Elm St
1	102	Jane Smith	jane@example.com	2345678901	456 Oak St
3	104	Alice Brown	alice@example.com	3456789012	789 Pine St
4	105	Bob Martin	bob@example.com	4567890123	321 Maple Ave

```
plt.figure(figsize=(8, 4))
df['Name'].value_counts().plot(kind='bar', color='blue')
```



```
plt.figure(figsize=(8, 4))
df_no_duplicates['Name'].value_counts().plot(kind='bar',
color='blue')
```



```
df_no_duplicates.to_csv("non_redundant_data.csv",
index=False)
print("File Saved")
```

Output:

```
▶ print(pd.read_csv("non_redundant_data.csv"))

[2]:   Customer_ID      Name        Email       Phone      Address
  0          101    John Doe  john@example.com  1234567890  123 Elm St
  1          102  Jane Smith  jane@example.com  2345678901  456 Oak St
  2          104  Alice Brown  alice@example.com  3456789012  789 Pine St
  3          105  Bob Martin  bob@example.com  4567890123  321 Maple Ave
```

Learning Outcomes:

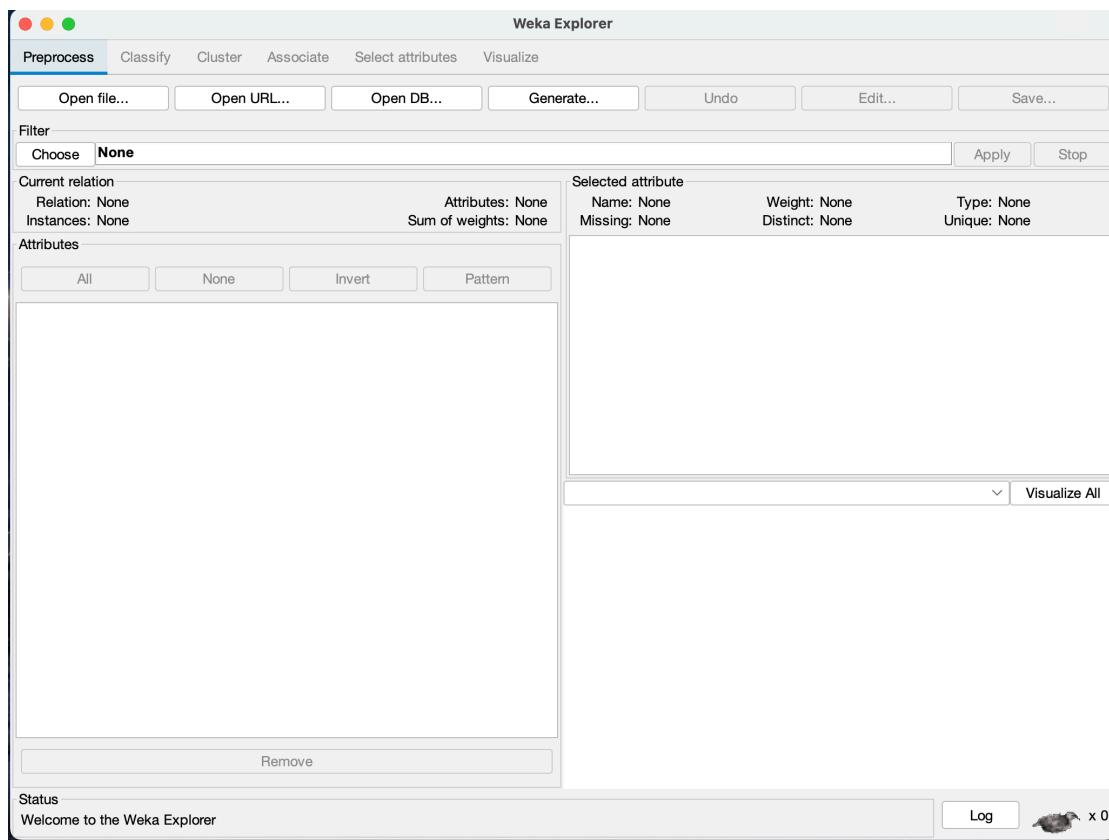
Experiment 4

Aim: Introduction to WEKA tool.

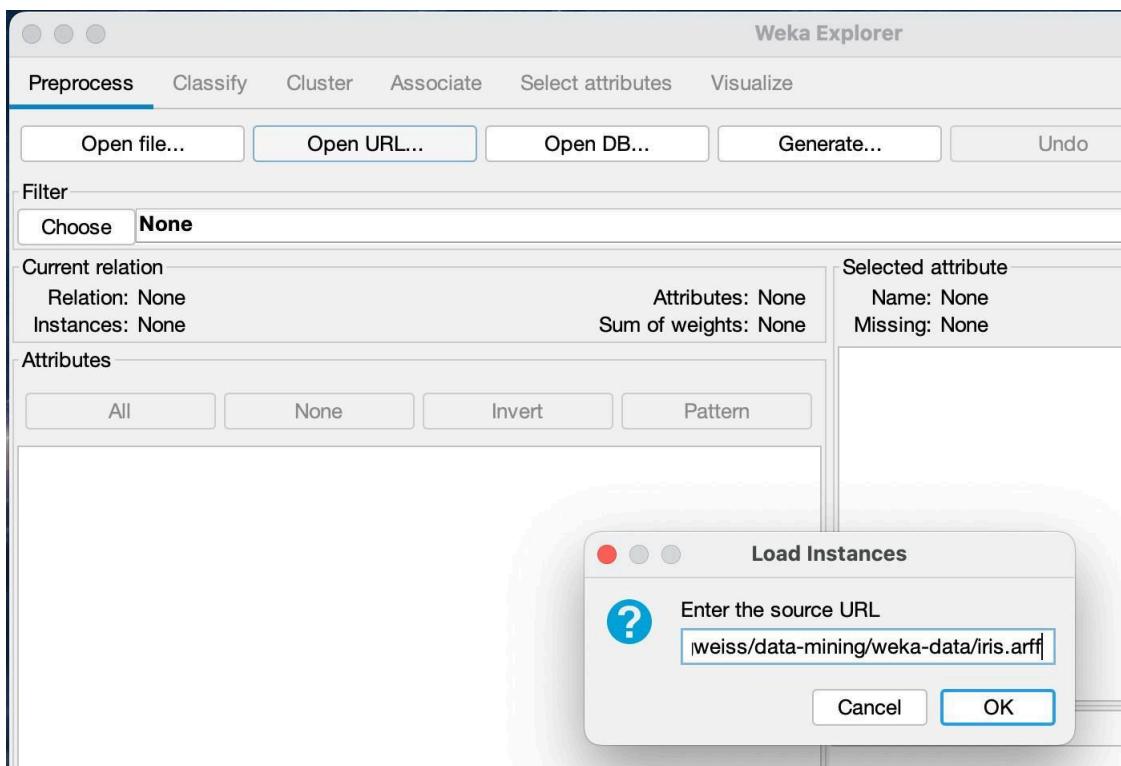
Theory:



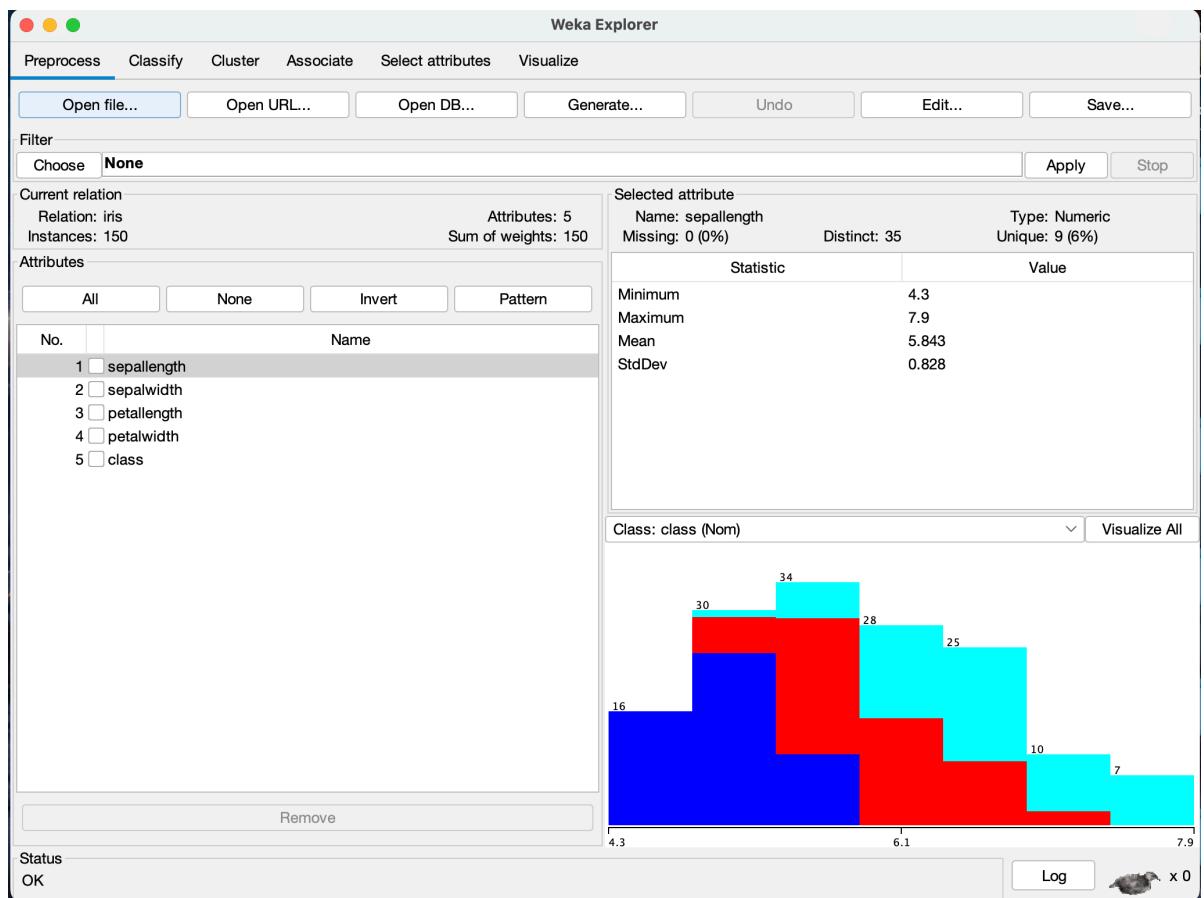
Homescreen of Weka



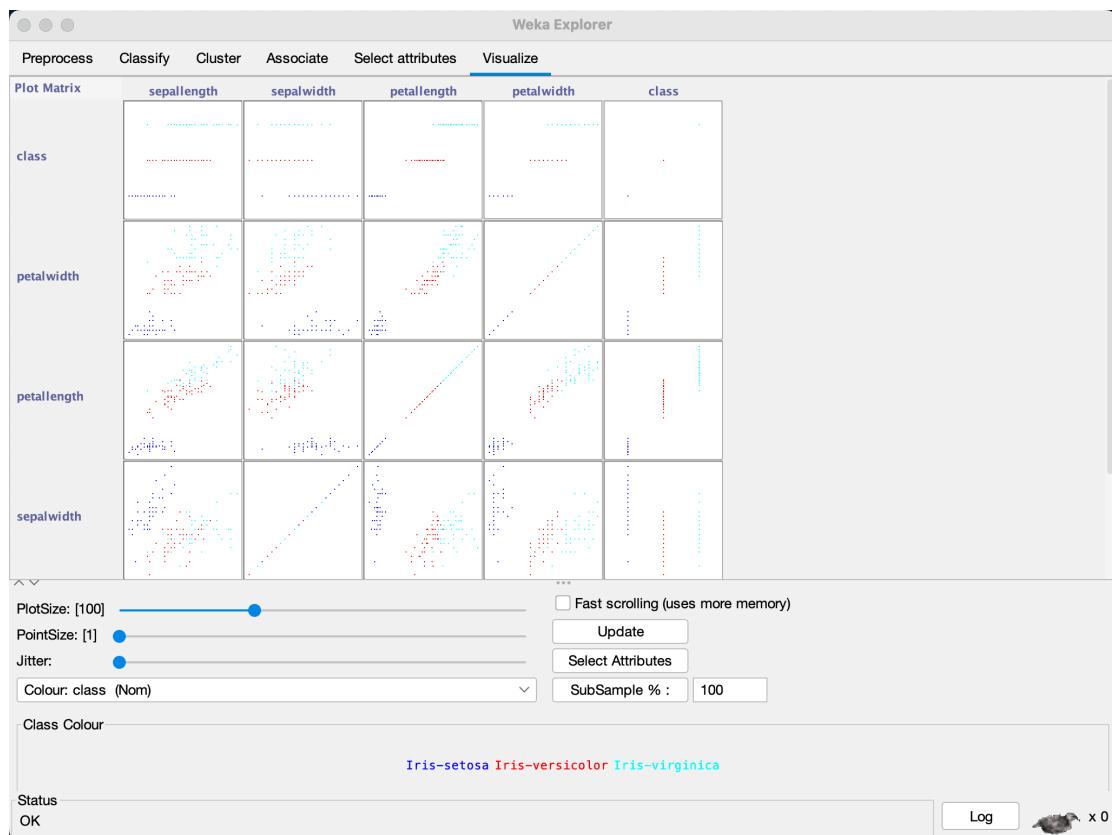
Weka Explorer



Opening Iris Dataset using URL



Preprocess Tab in Weka



Visualisation of Dataset

Learning Outcomes:

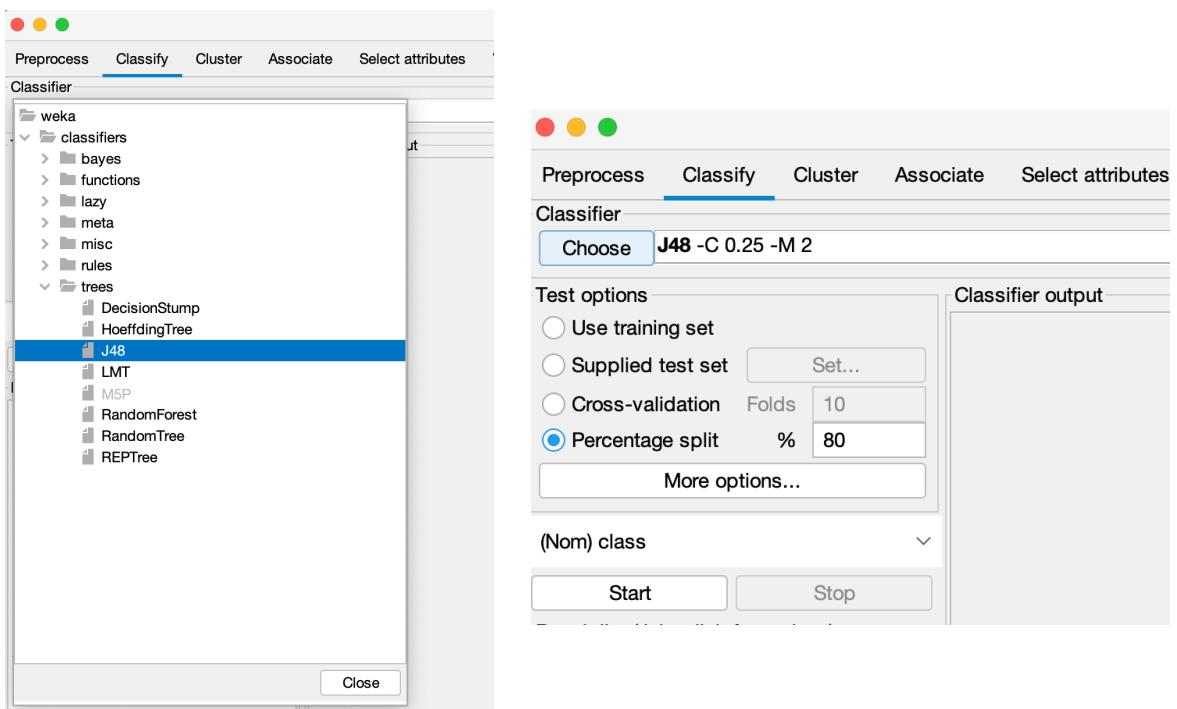
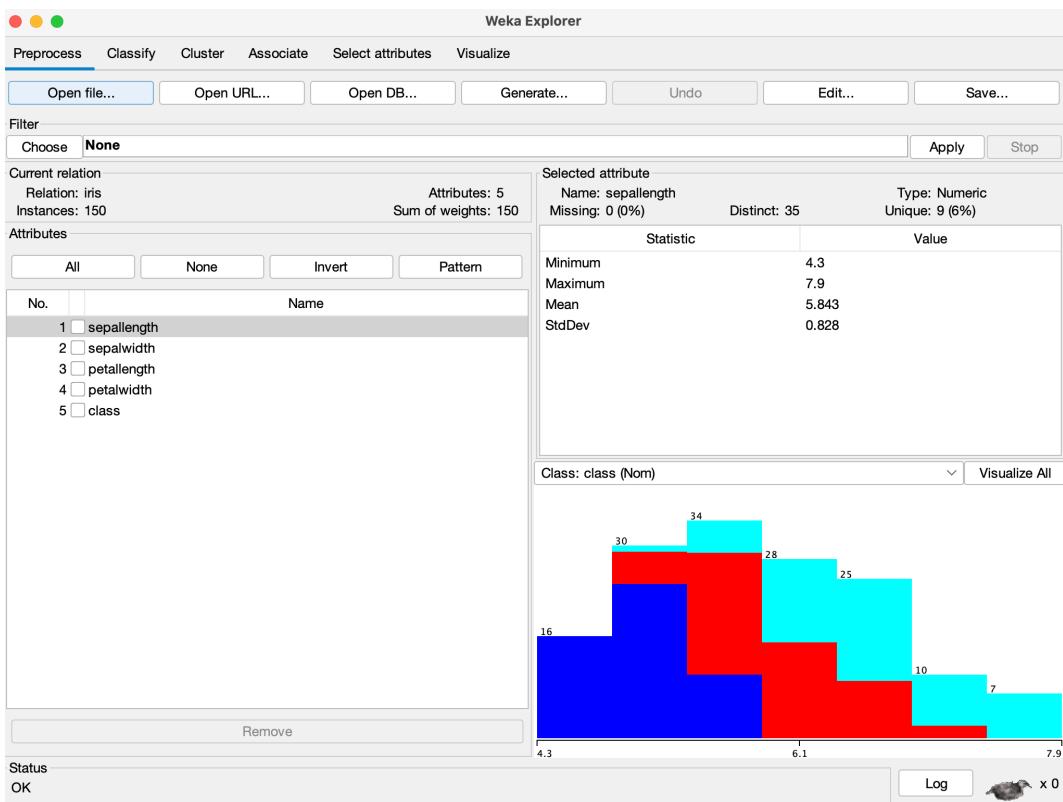
Experiment 5

Aim: Implementation of Classification technique on ARFF files using WEKA.

Theory:

Dataset:

Snapshots:



Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier Choose J48 -C 0.25 -M 2

Test options

- Use training set
- Supplied test set Set...
- Cross-validation Folds 10
- Percentage split % 80

More options...

(Nom) class

Start Stop

Result list (right-click for options)

01:33:34 - trees.J48

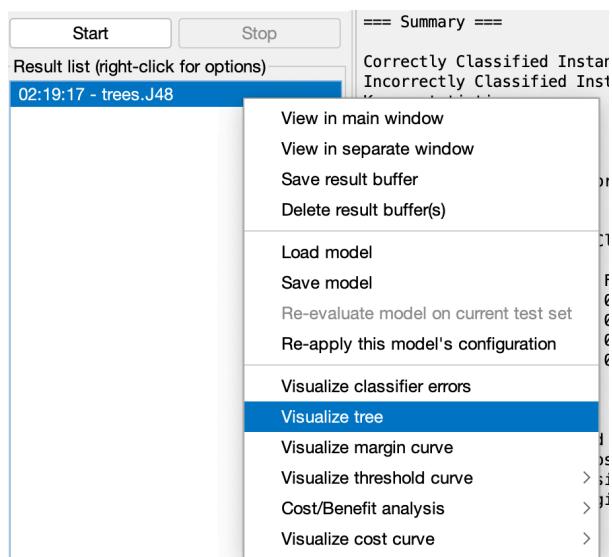
Classifier output

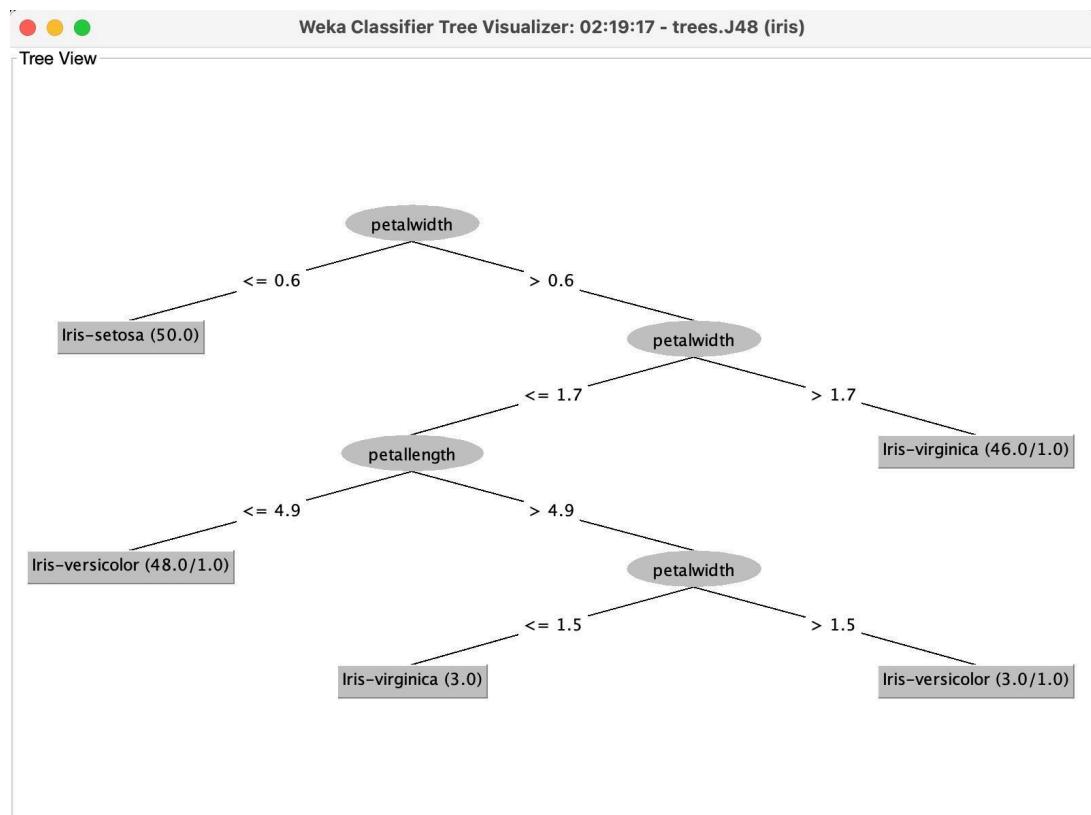
```
== Run information ==
Scheme: weka.classifiers.trees.J48 -C 0.25 -M 2
Relation: iris
Instances: 150
Attributes: 5
sepallength
sepalwidth
petallength
petalwidth
class
Test mode: split 80.0% train, remainder test
== Classifier model (full training set) ==
J48 pruned tree
-----
petalwidth <= 0.6: Iris-setosa (50.0)
petalwidth > 0.6
| petalwidth <= 1.7
| | petallength <= 4.9: Iris-versicolor (48.0/1.0)
| | petallength > 4.9
| | | petalwidth <= 1.5: Iris-virginica (3.0)
| | | petalwidth > 1.5: Iris-versicolor (3.0/1.0)
| | petalwidth > 1.7: Iris-virginica (46.0/1.0)

Number of Leaves : 5
Size of the tree : 9

Time taken to build model: 0.01 seconds
== Evaluation on test split ==
Time taken to test model on test split: 0 seconds
```

Status OK Log x 0





Output:

==== Run information ===

Scheme: weka.classifiers.trees.J48 -C 0.25 -M
 2 Relation: iris
 Instances: 150
 Attributes: 5
 sepallength
 sepalwidth
 petallength
 petalwidth
 class
 Test mode: split 80.0% train, remainder test

==== Classifier model (full training set)

==== J48 pruned tree

```
-----
petalwidth <= 0.6: Iris-setosa (50.0)
petalwidth > 0.6
|   petalwidth <= 1.7
|   |   petallength <= 4.9: Iris-versicolor (48.0/1.0)
|   |   petallength > 4.9
|   |   |   petalwidth <= 1.5: Iris-virginica (3.0)
|   |   |   petalwidth > 1.5: Iris-versicolor (3.0/1.0)
|   petalwidth > 1.7: Iris-virginica
```

(46.0/1.0) Number of Leaves : 5

Size of the tree : 9

Time taken to build model: 0.01 seconds

==== Evaluation on test split ===

Time taken to test model on test split: 0 seconds

==== Summary ===

Correctly Classified Instances	30	100	%
Incorrectly Classified Instances	0	0	%
Kappa statistic	1		
Mean absolute error	0.0105		
Root mean squared error	0.0166		
Relative absolute error	2.3665 %		
Root relative squared error	3.5274 %		
Total Number of Instances	30		

==== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	1.000	0.000	1.000	1.000	1.000	1.000	1.000	1.000	Iris-setosa
	1.000	0.000	1.000	1.000	1.000	1.000	1.000	1.000	Iris-versicolor
	1.000	0.000	1.000	1.000	1.000	1.000	1.000	1.000	Iris-virginica
Weighted Avg.	1.000	0.000	1.000	1.000	1.000	1.000	1.000	1.000	

==== Confusion Matrix ===

a	b	c	<-- classified as
11	0	0	a = Iris-setosa
0	10	0	b = Iris-versicolor
0	0	9	c = Iris-virginica

Learning Outcomes:

Experiment 6

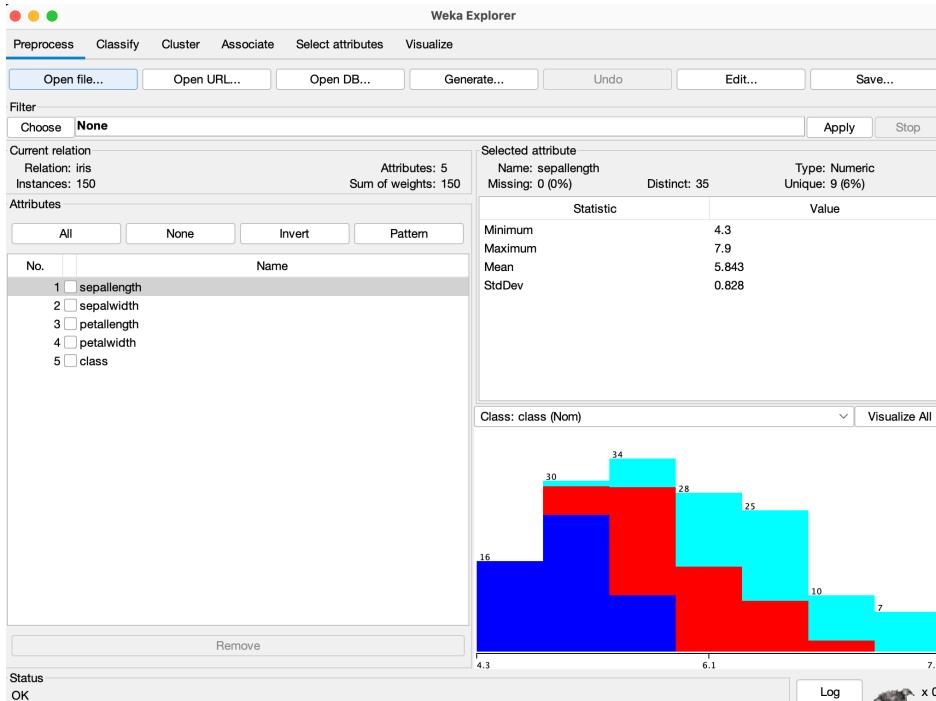
Aim: Implementation of Clustering technique on ARFF files using WEKA.

Theory:

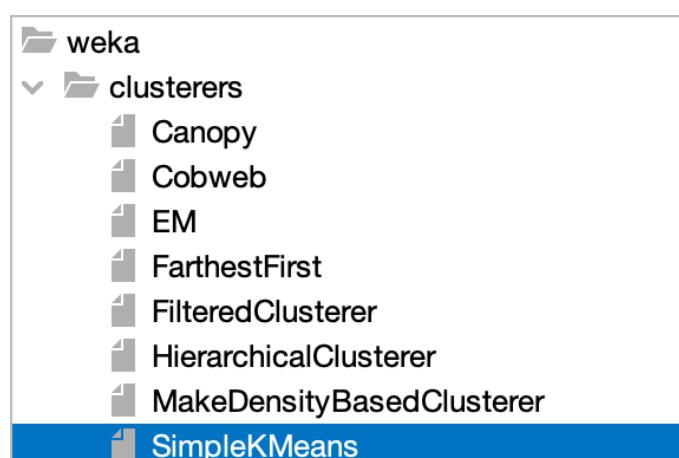
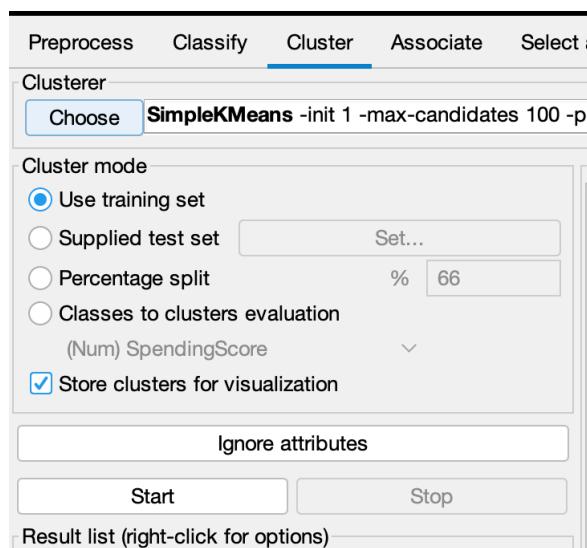
Dataset:

Procedure:

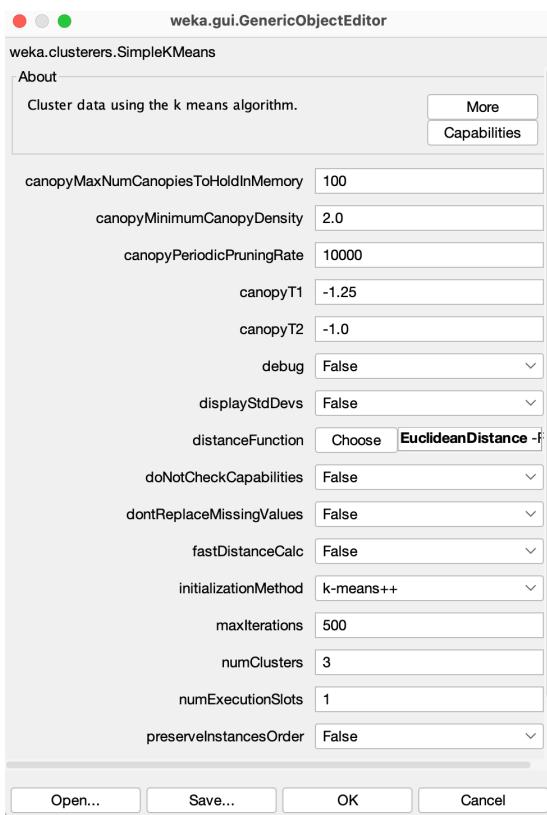
- WEKA GUI Chooser -> Explorer -> Open .arff file



- Go to Cluster tab -> Under Cluster Mode, keep: Use training set -> Click on Choose Pick: SimpleKMeans



3. Set K-Means Parameters: [numClusters = 3, Distance function: EuclideanDistance ,Initialization method: k-means++] -> Click OK



5. Run the clustering

Output:

```

Clusterer output
Instances: 150
Attributes: 5
sepallength
sepalwidth
petallength
petalwidth
class
Test mode: evaluate on training data

== Clustering model (full training set) ==

kMeans
=====

Number of iterations: 7
Within cluster sum of squared errors: 62.1436882815797

Initial starting points (random):
Cluster 0: 6.1,2.9,4.7,1.4,Iris-versicolor
Cluster 1: 6.2,2.9,4.3,1.3,Iris-versicolor

Missing values globally replaced with mean/mode

Final cluster centroids:
Attribute          Full Data      Cluster#
                  (150.0)       0           1
=====
sepallength        5.8433       6.262      5.006
sepalwidth         3.054        2.872      3.418
petallength        3.7587       4.906      1.464
petalwidth         1.1987       1.676      0.244
class              Iris-setosa   Iris-versicolor   Iris-setosa

Time taken to build model (full training data) : 0.01 seconds
== Model and evaluation on training set ==
Clustered Instances
0     100 ( 67%)
1      50 ( 33%)

```



Learning Outcomes:

Experiment 7

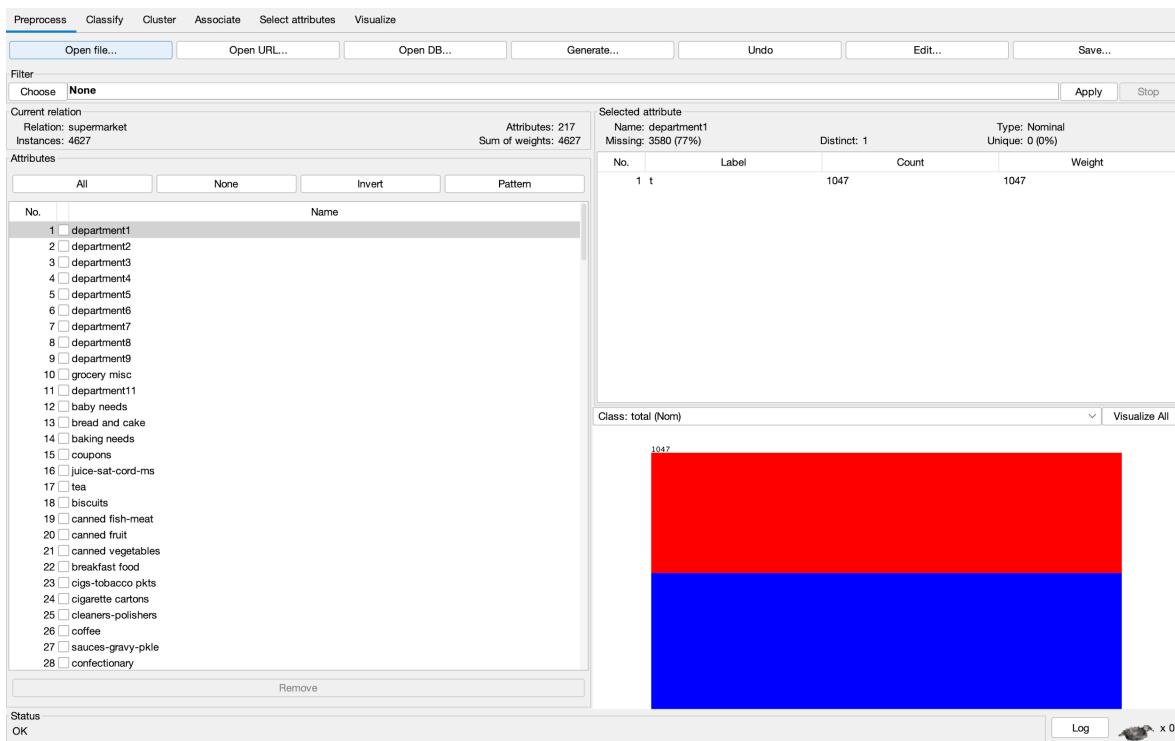
Aim: Implementation of Association Rule technique on ARFF files using WEKA.

Theory:

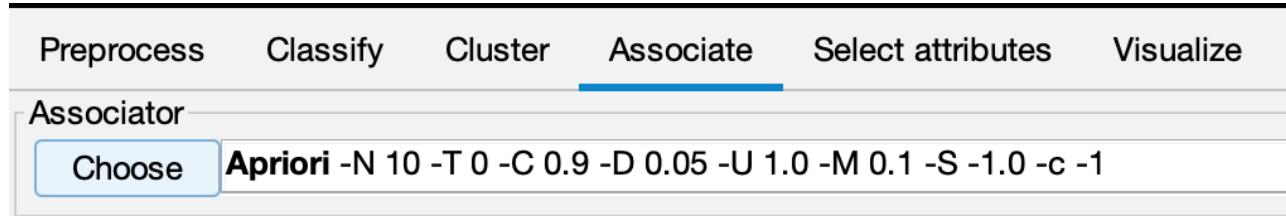
Dataset:

Procedure:

- WEKA GUI Chooser -> Explorer -> Open .arff file



- Go to Associate tab -> Choose Algorithm (Apriori)



- Run by clicking on Start

Output:

```

Associator output
== Run information ==
Scheme:      weka.associations.Apriori -N 10 -T 0 -C 0.9 -D 0.05 -U 1.0 -M 0.1 -S -1.0 -c -1
Relation:    supermarket
Instances:   4627
Attributes:  217
               [list of attributes omitted]
== Associator model (full training set) ==

Apriori
=====
Minimum support: 0.15 (694 instances)
Minimum metric <confidence>: 0.9
Number of cycles performed: 17

Generated sets of large itemsets:
Size of set of large itemsets L(1): 44
Size of set of large itemsets L(2): 380
Size of set of large itemsets L(3): 910
Size of set of large itemsets L(4): 633
Size of set of large itemsets L(5): 105
Size of set of large itemsets L(6): 1

Best rules found:
1. biscuits=t frozen foods=t fruit=t total=high 788 => bread and cake=t 723    <conf:(0.92)> lift:(1.27) lev:(0.03) [155] conv:(3.35)
2. baking needs=t biscuits=t fruit=t total=high 760 => bread and cake=t 696    <conf:(0.92)> lift:(1.27) lev:(0.03) [149] conv:(3.28)
3. baking needs=t frozen foods=t fruit=t total=high 770 => bread and cake=t 705    <conf:(0.92)> lift:(1.27) lev:(0.03) [150] conv:(3.27)
4. biscuits=t fruit=t vegetables=t total=high 815 => bread and cake=t 746    <conf:(0.92)> lift:(1.27) lev:(0.03) [159] conv:(3.26)
5. party snack foods=t fruit=t total=high 854 => bread and cake=t 779    <conf:(0.91)> lift:(1.27) lev:(0.04) [164] conv:(3.15)
6. biscuits=t frozen foods=t vegetables=t total=high 797 => bread and cake=t 725    <conf:(0.91)> lift:(1.26) lev:(0.03) [151] conv:(3.06)
7. baking needs=t biscuits=t vegetables=t total=high 772 => bread and cake=t 701    <conf:(0.91)> lift:(1.26) lev:(0.03) [145] conv:(3.01)
8. biscuits=t fruit=t total=high 954 => bread and cake=t 866    <conf:(0.91)> lift:(1.26) lev:(0.04) [179] conv:(3)
9. frozen foods=t fruit=t vegetables=t total=high 834 => bread and cake=t 757    <conf:(0.91)> lift:(1.26) lev:(0.03) [156] conv:(3)
10. frozen foods=t fruit=t total=high 969 => bread and cake=t 877    <conf:(0.91)> lift:(1.26) lev:(0.04) [179] conv:(2.92)

```

Learning Outcomes:

Experiment 8

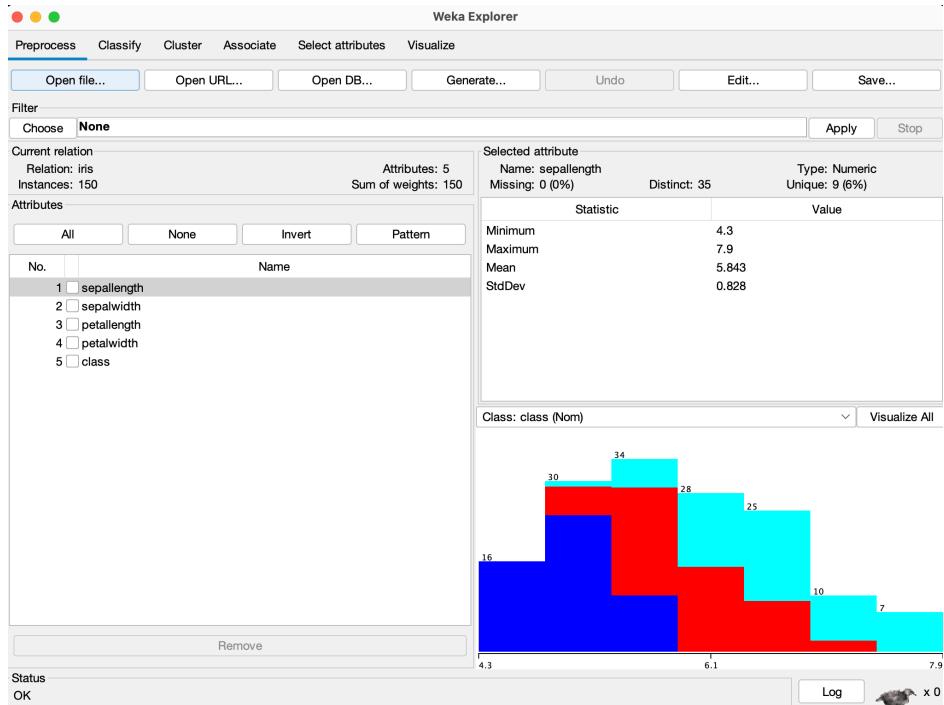
Aim: Implementation of Visualization technique on ARFF files using WEKA.

Theory:

Dataset:

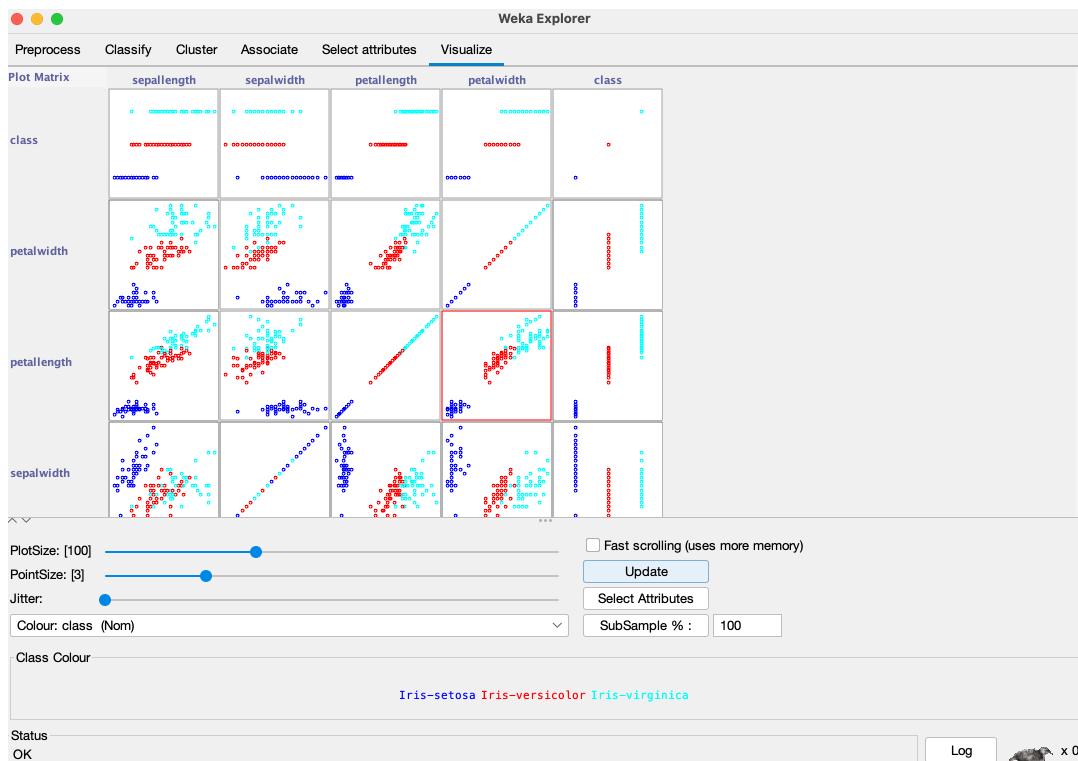
Procedure:

- WEKA GUI Chooser -> Explorer -> Open .arff file

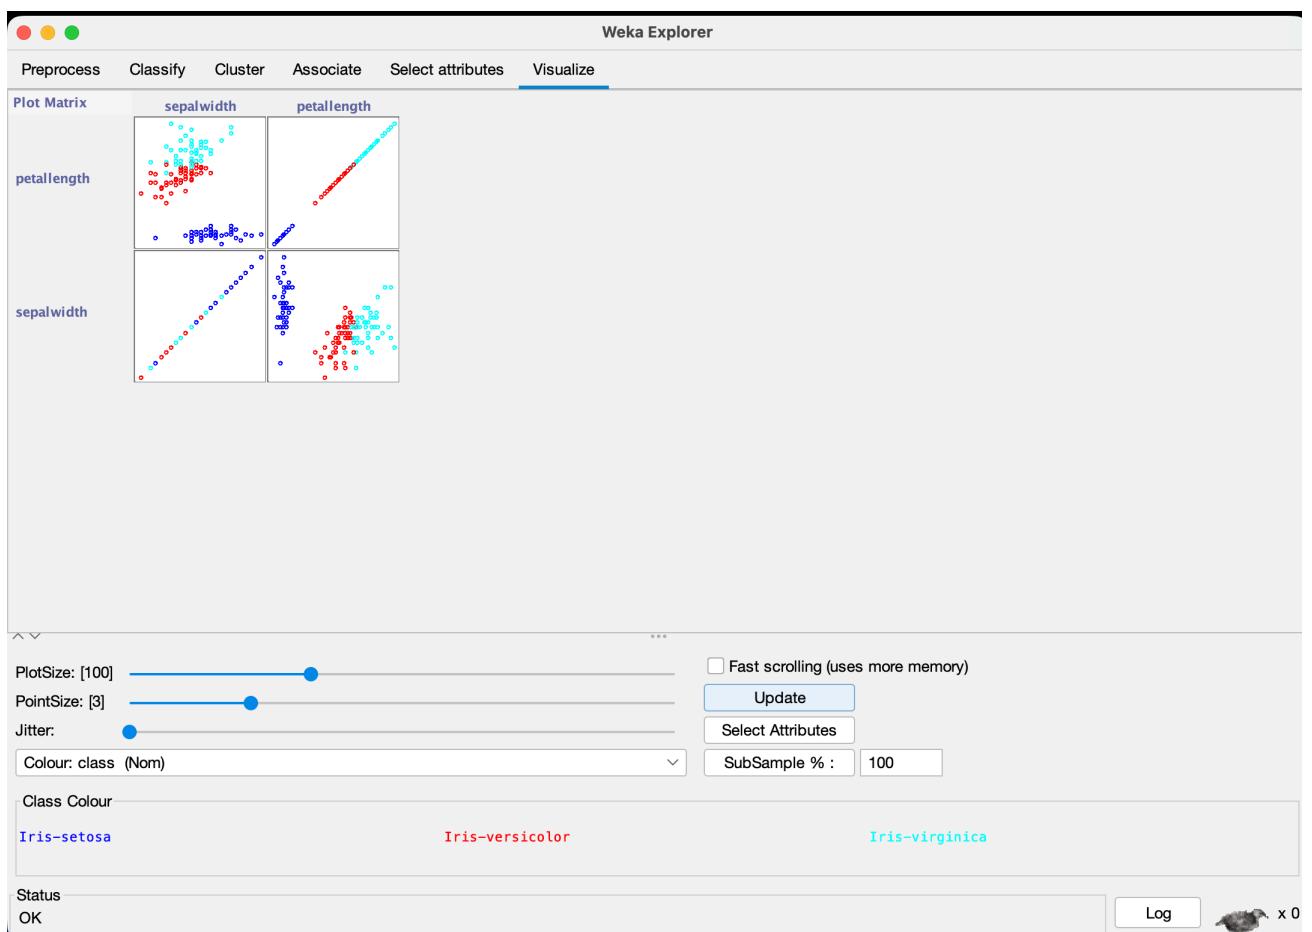
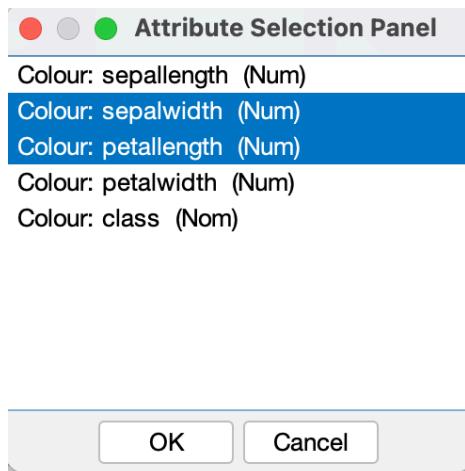


- Go to Visualize tab.

- Click Visualize All



b. Visualizing Selected Attributes



Learning Outcomes:

Experiment 9

Aim: Apply the Apriori Algorithm focusing on discriminating between patients with Parkinson's disease and other neurological disorders using the Voice Recording Dataset.

Theory:

Dataset:

Code:

```
import warnings, gc, io
warnings.filterwarnings('ignore')

!pip install pandas mlxtend

import pandas as pd

url = "https://archive.ics.uci.edu/ml/machine-learning-databases/parkinsons/parkinsons.data"
df = pd.read_csv(url)
df.head()

# Drop name column
df = df.drop("name", axis=1)

# Convert status column into labels
df['status'] = df['status'].map({0: "Healthy", 1: "Parkinsons"})

# Binning continuous features
for col in df.columns:
    if col != "status":
        df[col] = pd.qcut(df[col], q=3, labels=['Low', 'Medium', 'High'])

from mlxtend.preprocessing import TransactionEncoder

# Convert df to list of lists (transactions)
transactions = []
for i in range(len(df)):
    row_items = []
    for col in df.columns:
        row_items.append(f"{col}={df[col].iloc[i]}")
    transactions.append(row_items)

# Transaction Encoder
te = TransactionEncoder()
te_ary = te.fit(transactions).transform(transactions)

df_encoded = pd.DataFrame(te_ary, columns=te.columns_)
df_encoded.head()

from mlxtend.frequent_patterns import apriori, association_rules
# find frequent itemsets
freq_items = apriori(df_encoded, min_support=0.3,
use_colnames=True)
```

```

freq_items.sort_values(by="support", ascending=False).head()

rules = association_rules(freq_items, metric="confidence",
min_threshold=0.7)

# Filter rules predicting Parkinson's disease
rules_pd =
rules[rules['consequents'].astype(str).str.contains("status=Parkinsons")]

rules_pd.sort_values(by="confidence", ascending=False).head(10)

print("\nAssociation Rules Predicting Parkinson's Disease:\n")
display(rules_pd[['antecedents','consequents','support','confidence','lift']])

```

Output:

Association Rules Predicting Parkinson's Disease:						
	antecedents	consequents	support	confidence	lift	
0	(D2=High)	(status=Parkinsons)	0.307692	0.923077	1.224490	
1	(HNR=Low)	(status=Parkinsons)	0.302564	0.907692	1.204082	
8	(Jitter:DDP=High)	(status=Parkinsons)	0.312821	0.938462	1.244898	
26	(MDVP:APQ=High)	(status=Parkinsons)	0.328205	0.984615	1.306122	
28	(MDVP:Fhi(Hz)=Medium)	(status=Parkinsons)	0.312821	0.938462	1.244898	
...
560	(Shimmer:APQ3=High, Shimmer:DDA=High)	(MDVP:Shimmer=High, status=Parkinsons, Shimmer...)	0.307692	0.923077	2.903226	
561	(Shimmer:APQ5=High)	(MDVP:Shimmer=High, status=Parkinsons, Shimmer...)	0.307692	0.923077	3.000000	
562	(MDVP:Shimmer=High)	(status=Parkinsons, Shimmer:APQ5=High, Shimmer...)	0.307692	0.923077	2.950820	
563	(Shimmer:APQ3=High)	(MDVP:Shimmer=High, status=Parkinsons, Shimmer...)	0.307692	0.923077	3.000000	
564	(Shimmer:DDA=High)	(MDVP:Shimmer=High, status=Parkinsons, Shimmer...)	0.307692	0.923077	3.000000	

135 rows × 5 columns

Learning Outcomes:

Experiment 10

Aim: Implement the K-Means Clustering and K-Nearest Neighbors (KNN) algorithms on the Parkinson's Disease Voice Recording Dataset from the UCI Machine Learning Repository. Perform classification and clustering of patients based on their vocal features, analyze clustering results, evaluate classification accuracy, and compare the performance of both methods.

Theory:

Dataset:

Code:

```
# Install required libraries
!pip install scikit-learn pandas numpy matplotlib seaborn -q

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split,
cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (classification_report,
accuracy_score,
adjusted_rand_score,
silhouette_score
confusion_matrix
e,
confusion_matrix
x) # Added

from sklearn.decomposition import PCA # Added PCA
import warnings
warnings.filterwarnings('ignore')
print("\n Parkinson's Analysis: K-Means & KNN")

# LOAD & PREPROCESS DATA
try:
    # Load from direct URL
    url = "https://archive.ics.uci.edu/ml/machine-learning-
databases/parkinsons/parkinsons.data"
    df = pd.read_csv(url)

    y = df['status'].values # 1 = Parkinson's, 0 = Healthy
    X = df.drop(['name', 'status'], axis=1)
    print(f"Dataset loaded: {X.shape[0]} samples, {X.shape[1]} features")
except Exception as e:
    print(f"Error loading data: {e}")
    exit()

# Split data for supervised learning (KNN)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Feature Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X) # For K-Means (all data)
X_train_scaled = scaler.fit_transform(X_train) # For KNN (train)
X_test_scaled = scaler.transform(X_test) # For KNN (test)
```

```
print("Data split and standardized")
# Apply PCA for visualization (fit on all scaled data)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)
print("PCA applied for visualization")

# K-MEANS CLUSTERING (UNSUPERVISED)
print("\nK-Means Clustering (k=2)")

# Apply K-Means with K=2 (since we know there are 2 classes)
kmeans_2 = KMeans(n_clusters=2, random_state=42, n_init=10)
cluster_labels = kmeans_2.fit_predict(X_scaled)

# Clustering Evaluation Metrics
sil_score = silhouette_score(X_scaled, cluster_labels)
ari_score = adjusted_rand_score(y, cluster_labels) # Compare
clusters to true labels

print(f"Silhouette Score: {sil_score:.4f}")
print(f"Adjusted Rand Index (ARI): {ari_score:.4f} (vs true
labels)")

# K-NEAREST NEIGHBORS (SUPERVISED)
print("\nK-Nearest Neighbors (KNN)")

# Find optimal K for KNN using cross-validation
k_values = range(1, 20, 2)
cv_scores = []
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train_scaled, y_train, cv=5,
scoring='accuracy')
    cv_scores.append(scores.mean())

# Get the best K
optimal_k = k_values[np.argmax(cv_scores)]
print(f"Optimal K found: {optimal_k} (CV Accuracy:
{max(cv_scores):.4f})")

# Train KNN with optimal K
knn_optimal = KNeighborsClassifier(n_neighbors=optimal_k)
knn_optimal.fit(X_train_scaled, y_train)

# Predictions
y_pred_test = knn_optimal.predict(X_test_scaled)

# Evaluation
test_accuracy = accuracy_score(y_test, y_pred_test)

print(f"\nKNN Test Accuracy (k={optimal_k}): {test_accuracy:.4f}")
print(f"Classification Report (Test Set):")
print(classification_report(y_test, y_pred_test,
```

RUDRA SHARMA

CSE-B

10417702722

```
target_names=['Healthy', "Parkinson's"]))

# VISUALIZATION
print("\nVisualizations")

fig, axes = plt.subplots(2, 2, figsize=(14, 12))
fig.suptitle('K-Means vs. KNN Analysis on Parkinson\'s Data',
font-size=16)

# --- Plot 1: K-Means Clustering (PCA) ---
ax1 = axes[0, 0]
ax1.scatter(X_pca[cluster_labels == 0, 0], X_pca[cluster_labels == 0, 1], c='blue', label='Cluster 0', alpha=0.6)
ax1.scatter(X_pca[cluster_labels == 1, 0], X_pca[cluster_labels == 1, 1], c='red', label='Cluster 1', alpha=0.6)
ax1.set_title('K-Means Clustering (PCA Projection)')
ax1.set_xlabel('Principal Component 1')
ax1.set_ylabel('Principal Component 2')
ax1.legend()
ax1.grid(True, alpha=0.3)

# --- Plot 2: KNN 'K' Selection ---
ax2 = axes[0, 1]
ax2.plot(k_values, cv_scores, 'o-', c='purple')
ax2.axvline(optimal_k, color='red', linestyle='--',
label=f'Optimal K={optimal_k}')
ax2.set_title('KNN: Optimal K Selection')
ax2.set_xlabel('K (Number of Neighbors)')
ax2.set_ylabel('Cross-Validation Accuracy')
ax2.legend()
ax2.grid(True, alpha=0.3)

# --- Plot 3: KNN Confusion Matrix ---
ax3 = axes[1, 0]
cm = confusion_matrix(y_test, y_pred_test)
sns.heatmap(cm, annot=True, fmt='d', cmap='Greens', ax=ax3,
            xticklabels=['Healthy', "Parkinson's"],
            yticklabels=['Healthy', "Parkinson's"])
ax3.set_title(f'KNN Confusion Matrix (k={optimal_k})')
ax3.set_ylabel('True Label')
ax3.set_xlabel('Predicted Label')

# --- Plot 4: Performance Comparison ---
ax4 = axes[1, 1]
methods = ['KNN (Accuracy)', 'K-Means (ARI)']
scores = [test_accuracy, ari_score]
colors = ['#4CAF50', '#2196F3']
bars = ax4.bar(methods, scores, color=colors)
ax4.set_title('Model Performance Comparison')
ax4.set_ylabel('Score (0.0 to 1.0)')
ax4.set_ylim(0, 1.0)
```

```

# Add text labels
for i, bar in enumerate(bars):
    yval = bar.get_height()
    ax4.text(bar.get_x() + bar.get_width()/2.0, yval,
f'{scores[i]:.3f}', va='bottom')
ax4.grid(True, axis='y', alpha=0.3)

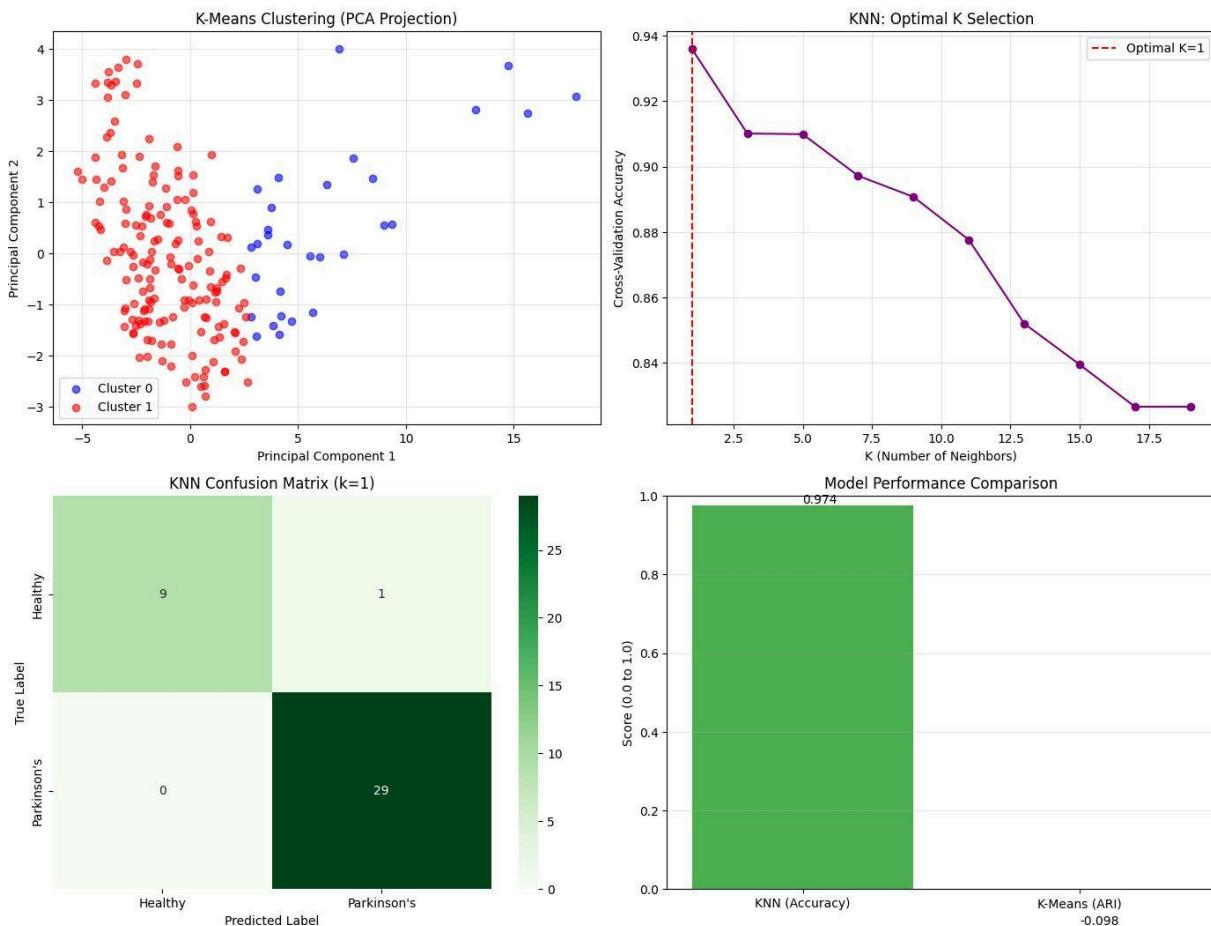
plt.tight_layout(rect=[0, 0.03, 1, 0.95]) # Adjust for suptitle
plt.savefig('kmeans_knn_minimal_analysis.png')
plt.show()

```

Output:

Classification Report (Test Set):				
	precision	recall	f1-score	support
Healthy	1.00	0.90	0.95	10
Parkinson's	0.97	1.00	0.98	29
accuracy			0.97	39
macro avg	0.98	0.95	0.97	39
weighted avg	0.98	0.97	0.97	39

K-Means vs. KNN Analysis on Parkinson's Data



RUDRA SHARMA

CSE-B

10417702722

Learning Outcomes: